

Inter Process Communication - 2

System V IPC

- Message Queues
- Shared Memory
- Semaphores

Common properties




- ✚ Each mechanism contains a table whose entries describe all instances of the mechanism.
- ✚ Each entry contains a numeric key, which is its user-chosen name.
- ✚ Each mechanism contains a “get” system call to create a new entry or to retrieve an existing one.
- ✚ Each IPC entry has a permissions structure that includes the user ID and group ID of the process that created the entry.
- ✚ Each mechanism contains a “control” system call to query status of an entry, to set status information, or to remove the entry from the system.

Message Queues

- ✚ A Message queue is a header pointing at a link list of messages.
- ✚ Each message contains a type field(4byte), followed by a data area.

Message Queues

Queue structure contains the following fields

-  Owners's uid and gid
-  Creator's uid and gid
-  Permissions

Message Queues

Queue structure contains the following fields

- + Owners's uid and gid
- + Creator's uid and gid
- + Permissions
- + Pointers to first and last messages on the linked list
- + Number of messages and total number of data bytes on the linked list

Message Queues

Queue structure contains the following fields

- + Owners's uid and gid
- + Creator's uid and gid
- + Permissions
- + Pointers to first and last messages on the linked list
- + Number of messages and total number of data bytes on the linked list
- + Max number of data bytes that can be on the linked list
- + Pid of last processes to send and receive messages
- + Timestamps – last msgsnd time, last msgrcv time, last change time

Message Queues


- ✚ The system call to get a message queue identifier is

`int msgget(key, flag)`
- ✚ It creates a new entry / retrieve a existing entry.
- ✚ Keys are numbers used to identify an IPC object on a UNIX system.
- ✚ Flag specifies the access permissions.
- ✚ Msgget returns a descriptor which is used as an index into an array of message queue headers.

Message Queues

- + New message queue is created
 - if key has value `IPC_PRIVATE`
 - If new key value is specified and `IPC_CREAT` flag is asserted in flags
- If `IPC_EXCL` and `IPC_CREAT` are both set and if message queue with the key value exists, then *msgget* fails

Message Queues



Message Send

- ✚ The system call to send a message is

```
int msgsnd(msgqid, msg, count, flag)
```

- *msgqid* - descriptor of message queue
- *msg* – pointer to the message structure

```
struct msgbuf
{
    long mtype;
    char mtext[1];
}
```

- *count* - size of the data array
- *flag* – action kernel should take if it runs out of internal buffer space

Message Send

Algorithm msgsnd

```
/*send a message*/
```

input : (1) message queue descriptor (2) address of message structure
(3) size of message (4) flags

output : number of bytes sent

$$\{$$

```
check legality of descriptor, permissions;
```

```
while(not enough space to store message)
```

$$\{$$

if(flags specify no to wait)

```
return;
```

```
sleep(until event enough space is available);
```

}

```
get message header;
```

```
read message text  from user space to kernel;
```

adjust data structures : enqueue message header, message header points to data, counts, timestamps, process ID;

```
wakeup all processes waiting to read message from queue;
```

}

Message Receive

✚ The system call to receive a message is

```
int msgrcv(msgqid, msg, count, type, flag)
```

- *msgqid* - descriptor of message queue
- *msg* – address of a user structure to contain received message
- *count* - size of the data array
- *type* - message type user wants to read
- *flag* – action kernel should do if no messages are on the queue

Message Receive

Algorithm msgrcv

/*receive message*/

input : (1) message queue descriptor (2) address of data array for incoming message

(3) size of data array (4) requested message type (5) flags

output : number of bytes in returned message

{

 check permissions;

loop:

 check legality of message descriptor;

 /* find message to return to user */

 if(requested message type == 0)

 consider first message on queue;

 else if(requested message type > 0)

 consider first message on queue with given type;

 else /* requested message type < 0 */

 consider first message of the lowest typed message on queue, such that its type is \leq absolute value of requested type;

Message Receive

```

If ( there is a message)
{
    adjust message size or return error if user size too small;
    copy message type, text from kernel space to user space;
    unlink message from queue;
    return;
}
/* no message */
if ( flags specify not to sleep )
    return with error;
sleep ( event message arrives on queue);
goto loop;
}
    
```

Message Control Operations

- ✚ A process uses the msgctl system call to query status and set parameters for the message queue

```
int msgctl( msgqid, cmd, buf )
```

- msgqid - identifies the message queue table entry
- Cmd - specifies the type of operation
- Buf - address of user buffer

Cmds available are

- IPC_STAT
 - Fetch the *msqid_ds* structure.

IPC_SET

- Set owner's user id, owners's group id, access modes, maximum number of bytes in the queue(only by super user).
- Permissions required by the process for the command
 - superuser privilege
 - process's effective uid equals creator's user id / owner's user id

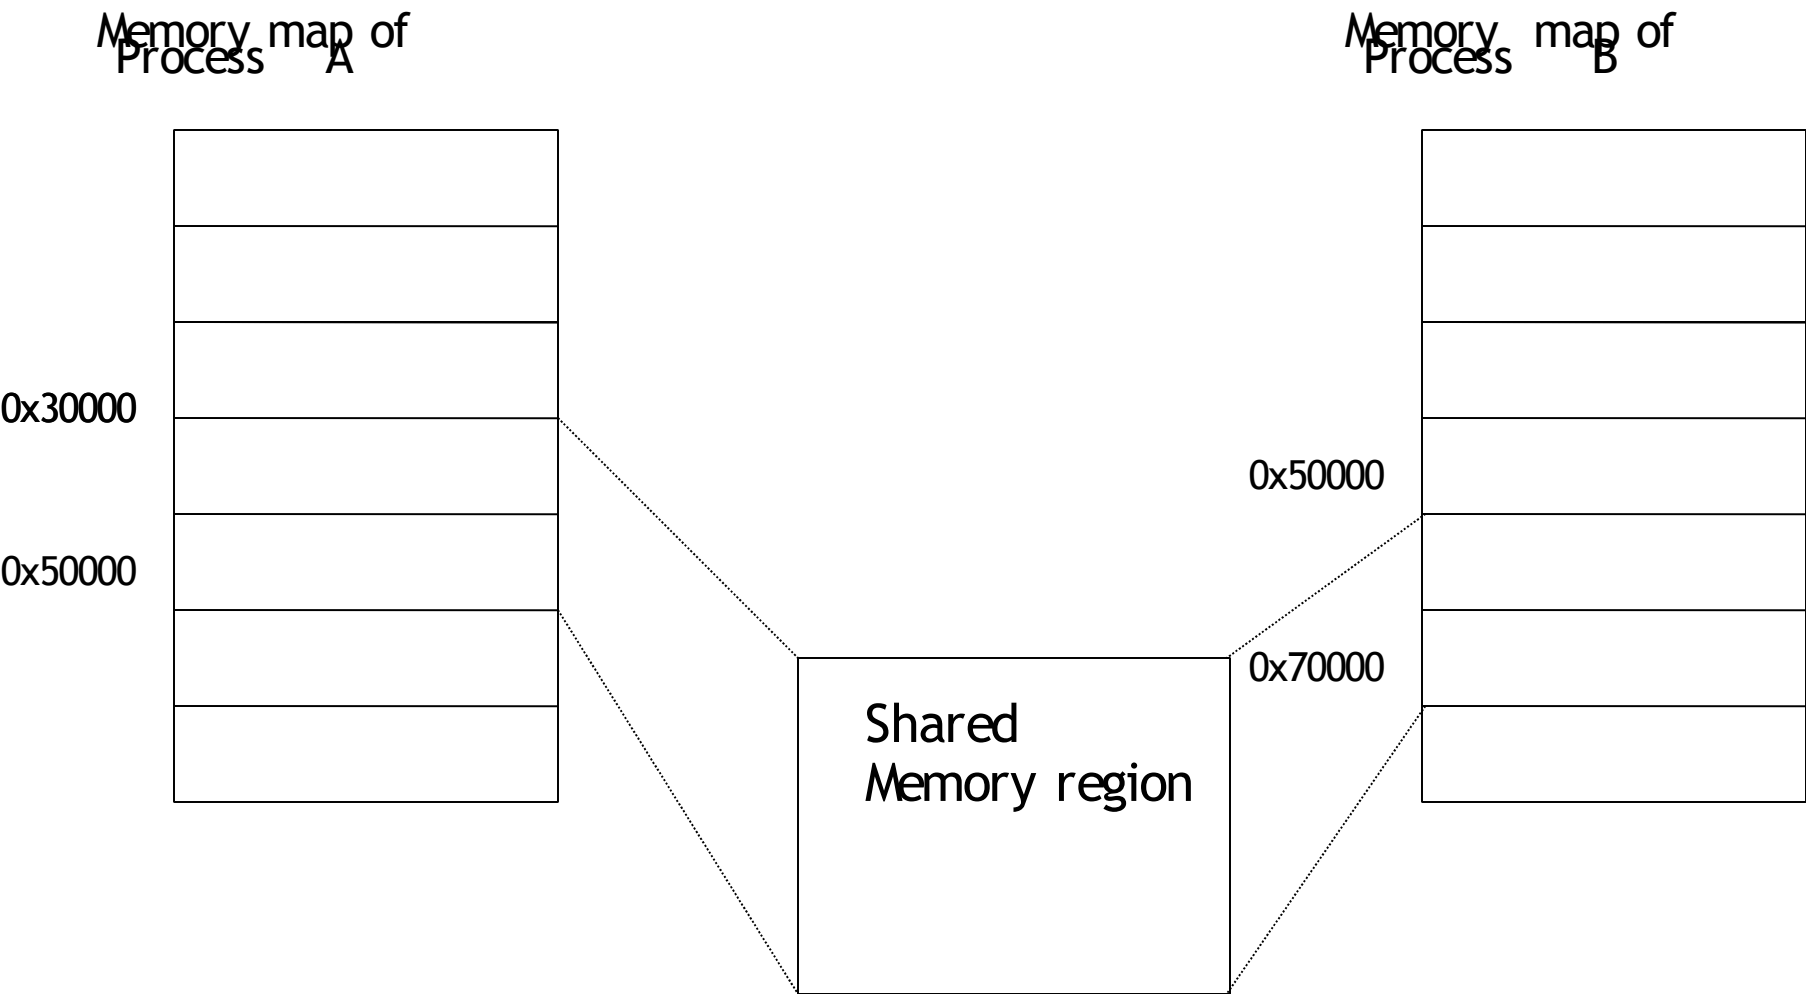
+ IPC_RMID

- Remove the message queue from the system.
- Removal is immediate.
- Permissions required by the process for the command
 - superuser privilege
 - process's effective uid equals creator's user id / owner's user id

Shared Memory




- + Shared memory is a portion of memory that is shared by multiple processes
- + Processes may attached this region to their virtual address space
- + Shared memory provides the fastest mechanism for processes to share data
- + If a process writes to a shared memory location, the new contents of that location are immediately visible to all processes sharing the region

Shared Memory



Shared Memory

Shared memory table contains the following fields

-  Owners's uid and gid
-  Creator's uid and gid
-  Permissions

Shared Memory

Shared memory table contains the following fields

- ✚ Owners's uid and gid
- ✚ Creator's uid and gid
- ✚ Permissions
- ✚ Size of the segment in bytes
- ✚ Pid of the creator and last operator
- ✚ Number of current attaches (nattach)
- ✚ Marked for deletion
- ✚ Timestamps – last attach time, last detach time, last change time

Shared Memory

- ✚ The system call to create a new region of shared memory or get an existing one is

```
int shmget(key, size, flag)
```

- ✚ It creates a new entry / retrieve a existing entry.
- ✚ Keys are numbers used to identify an IPC object on a UNIX system.
- ✚ Size is the number of bytes in the region
- ✚ Flag specifies the access permissions.
- ✚ shmget returns a segment descriptor

Shared Memory

Shared memory availability after the following system calls

- `Fork()` - After `fork()` child inherits the attached shared memory segments
- `Exec()` - After `exec()` all attached shared memory segments are detached (not destroyed)
- `Exit()` - On `exit()` all attached shared memory segments are detached (not destroyed)

Shared Memory Attach

- ✚ A Process attaches shared memory region to its virtual address space with the system call

$\text{virtaddr} = \text{shmat}(\text{shmid}, \text{addr}, \text{flags})$

- *Shmid* - descriptor returned by `shmget`
- *Addr* - virtual address where the user wants to attach the shared memory
- *Flag* - specifies whether the region is read-only and whether the kernel should round off the user specified address

Shared Memory Attach

Algorithm shmat /* attach shared memory*/

input : (1) shared memory descriptor

(2) virtual address to attach shared memory (3) flags

output : virtual address where shared memory was attached

```
{
    check validity of descriptor, permissions;
    if ( user specified virtual address )
    {
        round off virtual address, as specified by flags;
        check legality of virtual address, size of region;
    }
    else /* user wants kernel to find good address */
    {
        kernel picks virtual address : error if none available;
        attach region to process address space;
        if ( region being attached for the first time )
            allocate page tables, memory for region;
        return ( virtual address attached );
    }
}
```

Shared Memory Detach

- ✚ A Process detaches shared memory region from its virtual address space with the system call

int shmdt (addr)

- *Addr* - virtual address returned by shmat system call

Shared Memory Control

- ✚ A process uses the shmctl system call to query status and set parameters for the shared memory region

```
int shmctl( shmid, cmd, buf )
```

- Shmid - identifies the shared memory table entry
- Cmd - specifies the type of operation
- Buf - address of user buffer

Cmds available are

+ IPC_STAT

- Fetch the *shmid_ds* structure.

Cmds available are

IPC_SET

- Set owner's user id, owners's group id, access modes.
- Permissions required by the process for the command
 - superuser privilege
 - process's effective uid equals creator's user id / owner's user id

Cmds available are

 IPC_RMID

- Remove the shared memory segment set from the system if number of current attaches becomes 0.
- Permissions required by the process for the command
 - superuser privilege
 - process's effective uid equals creator's user id / owner's user id

Thank you